

# Finding Your Balance: Observability with Elastic and OpenTelemetry

A Hybrid Approach to Automatic and Manual  
Instrumentation in Python

# Justin Castilla

Senior Developer Advocate @ Elastic

Bluesky: @justincastilla.bsky.social

justin.castilla@elastic.co





# Talk Agenda:

- Observability Concepts
- OpenTelemetry Framework
  - Automatic Instrumentation
  - Manual Instrumentation
- A Case for a Hybrid Approach
- Hybrid Observability in Practice
- Questions Answered





# Observability Concepts

## What is Observability?

- Observability is the understanding of what is really happening to a system
- It's about the who, what, where, when, and why of your application's internal travel - not just state you are in at the moment
- Observability helps you determine what's happening inside your system by analyzing the data it outputs

# Observability Concepts

**Who** is using your application or visiting your site?

**What** can be done to fix detected anomalous behavior?

**Where** is the bottleneck in performance?

**When** is your service running the slowest?

**Why** is your site running slow in EMEA?

# Observability Concepts

## 3 Pillars of Observability:

### Logs

Records of events in a system, documenting operations, errors, and activities

### Traces

Tracks the path and interactions of a request through a system.

### Metrics

Quantitative measurements that track the performance and health of a system

We draw from all three individually and in combination to tell a story about your application.

# Observability Concepts

**Observability provides *context*.**

- Raw data alone *isn't* enough—context turns noise into actionable insights.
- Without context, it's like reading a list of random events with no understanding of how they're connected.
- Knowing a request failed is good; knowing why it failed is better.
- Observability is *not* about data volume but about providing meaningful answers through context.

# The OpenTelemetry Framework

## OpenTelemetry as an open standard

- A merger of two existing open source projects: OpenTracing and OpenCensus.
- An observability framework that is open source designed to work with any backend system.
- It provides standardized APIs, libraries, and tools to collect telemetry data.
- You may point your data from any supported platform to another with reasonably high confidence of 1:1 portability thanks to ECS (Elastic Common Schema)







# The OpenTelemetry Framework

## Instrumentation

*The process of adding observability features to your application to collect telemetry data, such as traces, metrics, and logs.*



# The OpenTelemetry Framework

## Automatic Instrumentation

- A script is invoked *before* your application is started and the service monitors activity of the available libraries. Supported modules can be found at the [OpenTelemetry Registry](#)
- This process typically covers HTTP, gRPC, database calls, and other common package interactions, depending on the instrumentation libraries enabled in your setup.





# The OpenTelemetry Framework

## Automagic is that easy!

- Very fast setup - most frameworks simply require that you run the wrapper script before your server startup
- Most common frameworks are supported (Django, Flask, FastAPI)
- Reference an options object or define environment variables containing sample frequency, packages to ignore, resource limits, etc.
- Demo: <http://localhost:4999>





# The OpenTelemetry Framework

## Automatic Challenges

- Lack of application-specific context
  - we only see the travel patterns, but not the data being sent
- Noise from irrelevant spans - this has the potential to create a LOT of unnecessary chatter
- Doesn't always align with your business context - creates uninformative, unactionable data





# The OpenTelemetry Framework



## Manual Instrumentation

- OTel code segments are written *into* your application to collect and transmit telemetry data.
- A collector receives telemetry data and sends them via the OTel protocol to a specific destination.
- You decide where, when, and what telemetry data is sent to the data platform, as well as how frequently and in what conditions.



# The OpenTelemetry Framework

## Manual Advantages

- Adds application-specific detail to traces
- Enables visibility into critical business workflows for your specific situation
- Low data transfer potential
- Filter attributes via processors.
- Demo: <http://localhost:5000>





# The OpenTelemetry Framework



## Manual Challenges

- Requires effort and discipline to implement
- Planning is necessary to write code to trace the paths and attributes that matter to you
- Risk of missing critical spans consequence of poor planning from the above
- Updates to your codebase may require updates to your manual implementation as well



# The OpenTelemetry Framework

## A hybrid approach

- Automatic provides a broad, sweeping view of your telemetry data as it traverses your application's landscape. You are able to see everything from one platform
- Manual allows you to select specific data to track that are of particular importance to your logic workflow. You choose the data to observe.





# The OpenTelemetry Framework



## A hybrid approach

By configuring your automatic instrumentation to *only* cover what paths and libraries you want covered and with a set interval rate, you can reduce noise and ensure high performance.

You can also explore logs, metrics, and traces *specific* to your observability needs



# Bringing it All Together

## Open Source for a reason

- You can use Elastic or **other** data platforms!
- Most modern platforms adhere to the OTel ECS standard
- It's an **observability** market, not a data platform market
- You can create your own instrumentation client for your own clients and libraries!

# Bringing it All Together

## Key Questions for a Balanced Approach

- What is your business-critical functionality?
  - What do you need to know, how often, and where does that info come from?
- What parts of your app already have rich telemetry?
  - Can you add to it?
  - How can you integrate that data into Open Telemetry?
- What's the overhead vs. value?
  - How important is this information to you temporally?
  - How much computation are you willing to allot to telemetry?

# Best Practices for Hybrid Instrumentation

## Strike your own balance

- Start with auto-instrumentation and layer manual spans incrementally.
  - This prevents duplication of efforts
- Evolve instrumentation as your application evolves.
  - Can you add to it?
  - Check in with your observability to identify any blindspots
- Use Elastic dashboards to validate and refine your telemetry strategy.
  - The data is there to be used!



# Hybrid Instrumentation for the win!

## Strike your own balance

- Observability *isn't* one-size-fits-all.
- Use automatic instrumentation for *quick* wins, and manual for *depth*.
- Elastic bridges raw data and actionable *insights*.
- Find your *balance* and focus on what matters most to your users.

# Hybrid Instrumentation for the win!

## Resources

- OpenTelemetry Demo ([Elastic fork](#))
- OpenTelemetry Python [SDK](#).
- Elastic Observability [docs](#).
- GitHub repo with [demo code](#).



# Thank you!

## Justin Castilla

Senior Developer Advocate @ Elastic

Bluesky: @justincastilla.bsky.social

justin.castilla@elastic.co

